

## **Part 1: Theoretical Analysis**

### **1. Short Answer Questions**

**Q1: How AI-driven code generation tools (e.g., GitHub Copilot) reduce development time. What are their limitations?**

How they reduce development time:

1. Boilerplate generation: They rapidly scaffold repetitive code structures, such as test setups, REST endpoints, and getters/setters.
2. Autocomplete and snippets: Offer context-sensitive code recommendations for multi-line blocks or in the middle of lines.
3. Help with documentation: They expedite clarity by drafting function documents and comments.
4. Error reduction: Reduces the likelihood of typographical errors by auto-completing common patterns.

Limitations:

1. Context understanding: Could recommend code for the project's domain that is syntactically correct but semantically incorrect.
2. Variance in quality: The recommendations still need developer review because they may be inadequate or unsafe.
3. Pattern dependency: Struggles with new or extremely specific logic; heavily relies on learned templates.
4. IP and bias: The possibility of recommending code that is similar to training examples that are protected by copyright.

**Q2: Compare supervised and unsupervised learning in the context of automated bug detection.**

Aspect	Supervised Learning	Unsupervised Learning
Data required	Labeled examples (buggy vs. non-buggy code)	Unlabeled data—code only
Detection method	Learns explicit bug patterns from labeled data	Learns normal behavior; flags outliers/anomalies
Accuracy & precision	Generally, more accurate when labeled data is abundant	More exploratory—useful for unknown bug types
Examples	Classifiers trained on past bug reports; vulnerability scans	Anomaly detection in runtime logs or unusual code commits
Limitations	Needs significant labeled data; brittle to novel bugs	Higher false positives; less precise without refinement

Q3: Why is bias mitigation critical when using AI for user experience personalization?

1. Fairness for all users: Steer clear of promoting preexisting biases, such as by displaying culturally biased content to particular groups of people.
2. Legal and ethical compliance: Businesses may face legal problems if they discriminate against people based on protected characteristics (such as gender, race, etc.).
3. Engagement and trust: Customers need to believe that the product treats them fairly; unjust personalization undermines loyalty and trust.
4. Long-term viability: Wider adoption and higher user satisfaction are the results of inclusive and balanced personalization.

## 2. Case Study Analysis

AIOps increases the effectiveness of software deployment by:

### 1. Automatic rollbacks and proactive anomaly detection

AI systems keep an eye on deployment metrics (such as error rates and response times); if they notice any irregularities, they can automatically roll back or send out alerts, which will cut down on downtime and mean time to recovery (MTTR).

### 2. Astute resource distribution and pipeline coordination

AI parallelizes deployment processes, anticipates bottlenecks, and optimizes resource provisioning across environments. For instance, allocating deployment queues according to risk assessment or dynamically scaling test environments during peak runs.

## **Task 1: Comparison & Analysis**

The objective of sorting a list of dictionaries by a given key is accomplished by both the manual and AI-generated code. The main distinction is in how flexibility and error handling are handled. Since `x[sort_key]` is used in the manual version, it is assumed that the target key is present in every dictionary. The AI-recommended version, on the other hand, employs `d.get(key_name, None)`, which prevents `KeyError` exceptions by offering a fallback (`None`) in the event that the key is missing.

The time complexity of both implementations is  $O(n \log n)$ , which is the same from an efficiency perspective because they both make use of Python's built-in `sorted()` function with a key extractor. Though this difference is insignificant in the majority of real-world scenarios, the AI-generated version might be marginally less performant in large datasets due to the overhead of the `.get()` method.

The AI version performs better in terms of functionality, particularly in situations where input data might include missing keys. Although it requires cleaner input, the manual version is easier to use and marginally faster.

In conclusion, manual code works best in controlled environments with guaranteed

structure, whereas AI-generated code is safer and more versatile for messy real-world data.

## **Task 2: Word Summary**

Software testing efficiency and coverage are greatly increased by AI-enhanced test automation tools such as Testim.io and Selenium IDE with AI plugins. We used both valid and invalid credentials in our automated login tests for this task. The AI plugins automatically managed dynamic wait times, proactively identified problematic tests, and recommended intelligent selectors (such as stable element locators based on semantic analysis).

By automatically repairing damaged selectors when UI elements change, AI tools lower maintenance overhead when compared to manual testing. Additionally, they identify timing or visual problems that conventional scripts might overlook. More significantly, AI learns user behavior patterns over time to help prioritize high-risk paths, allowing for better coverage with fewer test cases.

A 100% success rate was demonstrated by our test results: valid logins were accepted and invalid logins were appropriately rejected. Particularly in agile, quick-paced projects, AI testing tools improve the speed, scalability, and resilience of regression testing.

## **Part 3: Ethical Reflection**

Ethical Analysis: Predictive Resource Allocation's Bias and Fairness

Context: To model issue priority classification within an organization, the Task 3 predictive model—which was based on the Breast Cancer Dataset—has been modified. Although technically sound, when implemented in a real-world setting, ethical issues surface.

Potential Biases in the Dataset

1. Bias in Synthetic Labeling:

Task 3 involved the random generation of "priority" labels. Issue priorities may represent organizational or personal biases in a real-world context if they are determined by subjective human input (e.g., managers or team leads) (e.g., favoring certain departments or employees).

## 2. Groups Underrepresented:

The training data may underrepresent the problems of some teams if the model is applied to a company with teams of different sizes or historical representation. For instance, the model might learn to assign "low priority" to future issues from a small engineering team more frequently if critical bugs from that team are uncommon in the data.

## 3. Drift in Feature Relevance:

The features of the breast cancer dataset are medical in nature and have nothing to do with business concerns. The model may prioritize irrelevant signals, enhancing noise over fairness, if it is implemented without context-specific feature engineering.

## How Fairness Tools Can Help

An extensive open-source toolkit for identifying and reducing biases is IBM AI Fairness 360 (AIF360). Here's how it might be useful in this situation:

### 1. Identifying Bias:

Disparate impacts across sensitive attributes, such as gender, role, or team, can be evaluated by AIF360.

It can identify systemic benefits and drawbacks by comparing metrics such as True Positive Rate (TPR) or False Negative Rate (FNR) between groups.

### 2. Reducing Bias:

To balance representation, the toolkit offers pre-processing techniques (such as reweighting data).

To guarantee consistency across predictions, it provides post-processing methods like equalized odds adjustments and in-processing choices like adversarial debiasing.

### 3. Transparency and Accountability:

AIF360 guarantees that ethical performance is incorporated into model evaluation rather than being an afterthought by including fairness metrics in addition to accuracy and F1 score.